

Assignment 2: Multi-Variable Linear Regression

```
1 import numpy as np
2 import pandas as pd
3 from keras.datasets import boston_housing
4 import matplotlib.pyplot as plt
5 import seaborn as sns
```

```
1 (train_data, train_targets), (test_data, test_targets) = boston_housing.load_data()
2
3 print("Shape of training data: ", train_data.shape)
4 print("Shape of testing data: ", test_data.shape)
5 print("Shape of training targets: ", train_targets.shape)
6 print("Shape of testing targets: ", test_targets.shape)
7
8 feature_variances = np.var(train_data, axis=0)
9 max_var_idx, min_var_idx = np.argmax(feature_variances), np.argmin(feature_variances)
10
11 print("Index of feature with highest variance: ", max_var_idx, " with value:", feature_variances[max_var_idx])
12 print("Index of feature with lowest variance: ", min_var_idx, " with value:", feature_variances[min_var_idx])
```


Shape of training data: (404, 13)
Shape of testing data: (102, 13)
Shape of training targets: (404,)
Shape of testing targets: (102,)
Index of feature with highest variance: 9 with value: 27611.97237403192
Index of feature with lowest variance: 4 with value: 0.013723618009263812

```
1 mean = train_data.mean(axis=0)
2 std = train_data.std(axis=0)
3 train_data_norm = (train_data - mean) / std
4 test_data_norm = (test_data - mean) / std
```

```
1 # Add bias term
2 X_train = np.hstack([np.ones((train_data_norm.shape[0], 1)), train_data_norm])
3 X_test = np.hstack([np.ones((test_data_norm.shape[0], 1)), test_data_norm])
4 Y_train = train_targets.reshape(-1, 1)
5 Y_test = test_targets.reshape(-1, 1)
6
7 # Initialize parameters
8 np.random.seed(42)
9 W = np.random.randn(X_train.shape[1], 1)
10
11 # Hyperparameters
12 learning_rate = 0.1
13 momentum = 0.9
14 epochs = 200
15 velocity = np.zeros_like(W)
16
17 loss_history = []
18
19 # Gradient Descent
20 for epoch in range(epochs):
21     Y_pred = X_train @ W # Prediction
22     error = Y_pred - Y_train # Loss
23     loss = (1 / len(Y_train)) * np.sum(error**2) # MSE
24     loss_history.append(loss)
25     gradient = (2 / len(Y_train)) * X_train.T @ error # Gradient computation
26     velocity = momentum * velocity - learning_rate * gradient # Velocity update
27     W += velocity
28
29     if epoch % 25 == 0:
30         print(f"Epoch {epoch}: Loss = {loss:.4f}")
```

Epoch 0: Loss = 535.0718
Epoch 25: Loss = 36.0553
Epoch 50: Loss = 22.4927
Epoch 75: Loss = 22.1959
Epoch 100: Loss = 22.0093
Epoch 125: Loss = 22.0050
Epoch 150: Loss = 22.0049
Epoch 175: Loss = 22.0048

```
1 plt.figure(figsize=(12, 4))
2 plt.plot(loss_history)
3 plt.title("Loss Convergence")
4 plt.xlabel("Epochs")
5 plt.ylabel("Mean Squared Error")
6 plt.grid(True)
7 plt.show()
```



```
1 def plot_regression_line(X_raw, X_norm, feature_idx, feature_name):
2     X_feature_raw = X_raw[:, feature_idx]
3     plt.figure(figsize=(12, 5))
4     plt.scatter(X_feature_raw, Y_train, color="blue", alpha=0.5, label="True Values")
5     x_range_raw = np.linspace(X_feature_raw.min(), X_feature_raw.max(), 100)
6     x_range_norm = (x_range_raw - mean[feature_idx]) / std[feature_idx]
7
8     bias = w[0][0]
```

```
8     bias = w[0][0]
9     weight = W[feature_idx + 1][0]
10    y_pred = bias + weight * x_range_norm
11
12    plt.plot(x_range_raw, y_pred, color="red", label="Regression Line")
13    plt.title(f"Regression on Feature {feature_idx} ({feature_name})")
14    plt.xlabel(f"{feature_name} (Feature {feature_idx})")
15    plt.ylabel("House Price ($1000s)")
16    plt.legend()
17    plt.grid(True)
18    plt.show()
19    print('\n')
20
21    feature_names = [
22        "CRIM", "ZN", "INDUS", "CHAS", "NOX", "RM", "AGE",
23        "DIS", "RAD", "TAX", "PTRATIO", "B", "LSTAT"
24    ]
25
26    plot_regression_line(train_data, train_data_norm, max_var_idx, feature_names[max_var_idx])
27    plot_regression_line(train_data, train_data_norm, min_var_idx, feature_names[min_var_idx])
```

